

Homework #2 Solutions

February 25, 2021

Problem 1. Suppose we have an array $A[1 : n]$ of n *distinct* numbers. For any element $A[i]$, we define the **rank** of $A[i]$, denoted by $\text{rank}(A[i])$, as the number of elements in A that are strictly smaller than $A[i]$ plus one; so $\text{rank}(A[i])$ is also the correct position of $A[i]$ in the sorted order of A .

Suppose we have an algorithm **magic-pivot** that given any array $B[1 : m]$ (for any $m > 0$), returns an element $B[i]$ such that $m/3 \leq \text{rank}(B[i]) \leq 2m/3$ and has worst-case runtime $O(n)^1$.

Example: if $B = [1, 7, 6, 2, 13, 3, 5, 11, 8]$, then **magic-pivot**(B) will return one arbitrary number among $\{3, 5, 6, 7\}$ (since sorted order of B is $[1, 2, 3, 5, 6, 7, 8, 11, 13]$)

- (a) Use **magic-pivot** as a black-box to obtain a deterministic quick-sort algorithm with worst-case running time of $O(n \log n)$. (10 points)

Solution. A complete solution has three steps, algorithm, proof of correctness, and runtime analysis.

Algorithm: Recall that in quick sort, we pick the pivot p as any arbitrary index of the array A . In our modification, the only change is that we pick p as the index of output of **magic-pivot**; formally:

modified-quick-sort($A[1 : n]$):

- (a) If $n = 0$ or $n = 1$, return A .
- (b) Let $b = \text{magic-pivot}(A)$. Iterate over the array A and find the index p where $A[p] = b$.
- (c) Run **partition**(A, p) and let q be the index of the correct position of pivot.
- (d) Run **modified-quick-sort**($A[1 : q - 1]$) and **modified-quick-sort**($A[q + 1 : n]$).

Proof of correctness: By the correctness of **magic-pivot**, we will always be able to find an index p . Since original quick-sort works with any arbitrary choice of index p as pivot, the **modified-quick-sort** algorithm works correctly as well. (In fact, the entire point of using **magic-pivot** is to speedup quick-sort with minimal connection to its proof of correctness.)

Runtime analysis: Let $T(n)$ be the function for the worst-case runtime of the algorithm on inputs of length n . We claim that

$$T(n) \leq \max_{\frac{n}{3} \leq q \leq \frac{2n}{3}} (T(q-1) + T(n-q)) + O(n);$$

this is because by **magic-pivot** will return an element whose correct position in the array, which is the index q , will be between $n/3$ and $2n/3$; the rest is the same as quick-sort.

For any choice of $\alpha \in [\frac{1}{3}, \frac{2}{3}]$, consider the function $S_\alpha(n) = S_\alpha(\alpha \cdot n) + S_\alpha((1-\alpha) \cdot n) + O(n)$. By the definition of $T(n)$, we have $T(n) = O(\max_\alpha S_\alpha(n))$.

If we write the recursion tree for S_α , at every level the work done by the algorithm will be $C \cdot n$ (for some constant $C > 0$), and there will be at most $O(\max\{\log_{1/\alpha}(n), \log_{1/(1-\alpha)}(n)\})$ levels in the tree (this is similar to several other recursion trees we have written and so we omit it here.) This means that $S_\alpha(n) = O(n \cdot \max\{\log_{1/\alpha}(n), \log_{1/(1-\alpha)}(n)\})$. Finally, since $\alpha \in [\frac{1}{3}, \frac{2}{3}]$, $\max\{\log_{1/\alpha}(n), \log_{1/(1-\alpha)}(n)\} \leq \log_{3/2}(n) = \Theta(\log n)$ over all choices of α .

¹Such an algorithm indeed exists, but its description is rather complicated and not relevant to us in this problem.

Finally, this means that the runtime of the algorithm is $T(n) = O(n \log n)$, as desired.

Note: The above runtime analysis was a very formal way of proving the upper bound on $T(n)$ without making any assumptions. If, in your homework, you have simply stated that the “worst-case” of the recursion is when the split is most unbalanced, i.e., $T(n) \leq T(n/3) + T(2n/3) + O(n)$, which implies $T(n) = O(n \log n)$, you will receive the full grade. The goal of showing the proof in full generality was to also prove the unbalanced split case is indeed the worst case.

- (b) Use **magic-pivot** as a black-box to design an algorithm that given the array A and any integer $1 \leq r \leq n$, finds the element in A that has rank r in $O(n)$ time². (15 points)

Hint: Suppose we run **partition** subroutine in quick sort with pivot p and it places it in position q . Then, if $r < q$, we only need to look for the answer in the subarray $A[1 : q]$ and if $r > q$, we need to look for it in the subarray $A[q + 1 : n]$ (although, what is the new rank we should look for now?).

Solution. A complete solution has three steps, algorithm, proof of correctness, and runtime analysis.

Algorithm: The algorithm is as follows:

find-rank($A[1 : n], r$):

- (a) If $n = 1$, return $A[1]$.
- (b) Let $b = \mathbf{magic-pivot}(A)$. Iterate over the array A and find the index p where $A[p] = b$.
- (c) Run **partition**(A, p) and let q be the index of the correct position of pivot.
- (d) If $q = r$, return $A[q]$.
- (e) Else, if $q > r$, return **find-rank**($A[1 : q - 1], r$); otherwise, return **find-rank**($A[q + 1 : n], r - q$) (note the change in the value of second argument).

Proof of Correctness: Proof is by induction: our hypothesis is that **find-rank**(A, r) outputs the correct answer for any choice of n and $1 \leq r \leq n$.

The base case is true when $n = 1$, since in this case $r = 1$ and the element of rank 1 is $A[1]$.

For the induction step, suppose this is true for all choices of $n \leq i + 1$ and we prove it for $n = i + 1$. By the correctness of **magic-pivot** and **partition**, we know that q is the correct position of $A[q]$ in the sorted array after the partitioning step; in other words, rank of $A[q]$ is q .

So if $q = r$, outputting $A[q] = A[r]$ is the correct answer.

If $q > r$, this means that the element with rank r belongs to the sub-array $A[1 : q - 1]$ as these are the elements smaller than $A[q]$ and since $r < q$, $A[r] < A[q]$ also by definition of rank. Thus, by induction hypothesis, **find-rank**($A[1 : q - 1], r$) finds the element of rank r in $A[1 : q - 1]$ which is also the element of rank r in A , making the answer correct.

Finally, if $q < r$, the element of rank r belongs to $A[q + 1 : n]$. Note however since we are removing q elements with value smaller than $A[r]$ from consideration, when looking at $A[q + 1 : n]$, the element of rank r in A will have rank $q - r$ in $A[q + 1 : n]$. By induction hypothesis, **find-rank**($A[q + 1 : n], q - r$) will find this element, finalizing the proof.

Runtime analysis: Define $T(n)$ as the worst-case runtime of **find-rank** on any array of length n (and for any choice of r). We have

$$T(n) \leq \max_{\frac{n}{3} \leq q \leq \frac{2n}{3}} T(q) + O(n);$$

²Note that an algorithm with runtime $O(n \log n)$ follows immediately from part (a)—sort the array and return the element at position r . The goal however is to obtain an algorithm with runtime $O(n)$.

this is by exactly the same argument as in part (a). Given that $T(n)$ is a monotone function of n (runtime of algorithm on a larger input can only become larger), we have $T(n) \leq T(\frac{2n}{3}) + O(n)$. This means (by replacing $O(n)$ with $C \cdot n$ for some constant $C > 0$),

$$T(n) \leq T(2n/3) + C \cdot n \leq T(4n/9) + C \cdot (n + 2n/3) \leq C \cdot n \cdot \sum_{i=0}^{+\infty} (2/3)^i = O(n),$$

as the sum of a geometric series with ratio less than 1 converges to $O(1)$. As such, the runtime of **find-rank** is $O(n)$ as desired.

Problem 2. Suppose we have an array $A[1 : n]$ which consists of numbers $\{1, \dots, n\}$ written in some arbitrary order (this means that A is a *permutation* of the set $\{1, \dots, n\}$). Our goal in this problem is to design a very fast randomized algorithm that can find an index i in this array such that $A[i] \bmod 3 = 0$, i.e., $A[i]$ is divisible by 3. For simplicity, in the following, we assume that n itself is a multiple of 3 and is at least 3 (so a correct answer always exist). So for instance, if $n = 6$ and the array is $A = [2, 5, 4, 6, 3, 1]$, we want to output either of indices 4 or 5.

- (a) Suppose we sample an index i from $\{1, \dots, n\}$ uniformly at random. What is the probability that i is a correct answer, i.e., $A[i] \bmod 3 = 0$? (5 points)

Solution. There are exactly $n/3$ numbers in $\{1, \dots, n\}$ that are multiples of 3 (as n itself is a multiple of 3 there is no corner case). Since we are picking i uniformly at random, the probability that i is any of these numbers is exactly $(n/3)/n = 1/3$. So the answer is $1/3$.

- (b) Suppose we sample m indices from $\{1, \dots, n\}$ uniformly at random and with repetition. What is the probability that none of these indices is a correct answer? (5 points)

Solution. By part (a), the probability that each index is *not* correct is $1 - 1/3 = 2/3$. Since we are sampling each index *independently* (as it is with repetition), the probability that no index is correct among m trials is $(2/3)^m$.

Now, consider the following simple algorithm for this problem:

Find-Index-1($A[1 : n]$):

- Let $i = 1$. While $A[i] \bmod 3 \neq 0$, sample $i \in \{1, \dots, n\}$ uniformly at random. Output i .

The proof of correctness of this algorithm is straightforward and we skip it in this question.

- (c) What is the **expected** worst-case running time of **Find-Index-1**($A[1 : n]$)? Remember to prove your answer formally. (7 points)

Solution. Define a random variable $X \in [1 : +\infty]$ where $X = j$ if the number of times we run the while-loop is j (it is a random variable depending on the randomness of the algorithm). Each run of the algorithm takes $O(X)$ time (but this is a random variable and so we need to turn it into a formula); thus the expected worst-case runtime of the algorithm is $O(\mathbf{E}[X])$. So, we only need to compute $\mathbf{E}[X]$.

We have,

$$\begin{aligned}
 \Pr(X = j) &= \Pr(\text{first } j - 1 \text{ trials fail and } j\text{-th trial succeeds}) && \text{(by the definition of while-loop)} \\
 &= \Pr(\text{first } j - 1 \text{ trials fail}) \cdot \Pr(j\text{-th trial succeeds}) \\
 & && \text{(by independence of trials in different iterations)} \\
 &= (2/3)^{j-1} \cdot (1/3) && \text{(by part (b) and part (a), respectively)} \\
 &< (2/3)^j.
 \end{aligned}$$

As such, by the definition of expectation,

$$\mathbf{E}[X] = \sum_{j=1}^{\infty} \Pr(X = j) \cdot j \leq \sum_{j=1}^{\infty} (2/3)^{j-1} \cdot j = 9,$$

as the series converges to 9. So $O(\mathbf{E}[X]) = O(1)$, meaning that the expected worst-case runtime of the algorithm is $O(1)$.

The problem with **Find-Index-1** is that in the worst-case (and not in expectation), it may actually never terminate! For this reason, let us consider a simple variation of this algorithm as follows.

Find-Index-2($A[1 : n]$):

- For $j = 1$ to n :
 - Sample $i \in \{1, \dots, n\}$ uniformly at random and if $A[i] \bmod 3 = 0$, output i and terminate; otherwise, continue.
- If the for-loop never terminated, go over the array A one element at a time to find an index i with $A[i] \bmod 3 = 0$ and output it as the answer.

Again, we skip the proof of correctness of this algorithm.

- (d) What is the **worst-case running time** of **Find-Index-2**($A[1 : n]$)? What about its **expected** worst-case running time? Remember to prove your answer formally.

(8 points)

Solution. The worst-case runtime of the new algorithm happens when we finish the for-loop without success and then do a linear search over the array; both of these takes $\Theta(n)$ time so the worst-case runtime is $\Theta(n)$.

For the expected worst-case runtime, let us define two variables. We use $X \in \{1, \dots, n, n+1\}$ to denote the number of iterations of the first for-loop where $X = n+1$ means that the for-loop failed. So, when $X \leq n$, the runtime of the algorithm is $O(X)$ and when $X = n+1$, the runtime of the algorithm is $O(n)$ (for the first for-loop) plus another $O(n)$ (for the second for-loop); either way, the runtime of the algorithm is $O(X)$. We thus need to compute expected value of X to get the expected worst-case runtime of the algorithm.

$$\mathbf{E}[X] = \sum_{j=1}^{n+1} \Pr(X = j) \cdot j \leq \sum_{j=1}^{\infty} (2/3)^{j-1} \cdot j = 9,$$

where the calculations is exactly as in part (a). Thus, in this case also, the expected worst-case runtime of the algorithm is $O(1)$.

Problem 3. Given an array $A[1 : n]$ of a combination of n positive and negative integers, our goal is to find whether there is a sub-array $A[l : r]$ such that

$$\sum_{i=l}^r A[i] = 0.$$

Example. Given $A = [13, 1, 2, 3, -4, -7, 2, 3, 8, 9]$, the elements in $A[2 : 8]$ add up to zero. Thus, in this case, your algorithm should output *Yes*. On the other hand, if the input array is $A = [3, 2, 6, -7, -20, 2, 4]$, then no sub-array of A adds up to zero and thus your algorithm should output *No*.

Hint: Observe that if $\sum_{i=l}^r A[i] = 0$, then $\sum_{i=1}^{l-1} A[i] = \sum_{i=1}^r A[i]$; this may come handy!

- (a) Suppose we are promised that every entry of the array belongs to the range $\{-5, -4, \dots, 0, \dots, 4, 5\}$. Design an algorithm for this problem with worst-case runtime of $O(n)$. **(15 points)**

Hint: Counting sort can also be used to efficiently sort arrays with negative entries whose absolute value is not too large; we just need to “shift” the values appropriately.

Solution. A complete solution has three steps, algorithm, proof of correctness, and runtime analysis.

Algorithm: We start by constructing a prefix sum array B as follows.

- (a) $B[0] = 0, B[1] = A[1]$.
- (b) For $i = 2$ to n , $B[i] = B[i - 1] + A[i]$

This way $B[i] = \sum_{j=1}^i A[j]$. We know each element in the array is at most 5 and at least -5 , so we have that every $-5n \leq B[i] \leq 5n$.

We now design our algorithm for this part.

- (a) Create the prefix sum array B as above.
- (b) Initialize array C of size $10n + 1$ to be zero.
- (c) For $i = 0$ to n , $C[B[i] + 5n + 1] = C[B[i] + 5n + 1] + 1$.
- (d) For $i = 1$ to $10n + 1$, if $C[i] > 1$, return *Yes*.
- (e) Return *No*.

Proof of Correctness: The fact that for each i , $B[i] = \sum_{j=1}^i A[j]$, can be proven using induction on i (it is actually so simple that you do not need to provide a proof for it). The base case is when $i = 1$, $B[1] = A[1]$ by construction of B . We assume for some k $B[k] = \sum_{j=1}^k A[j]$. For $k + 1$, we know $B[k + 1] = B[k] + A[k + 1] = \sum_{j=1}^{k+1} A[j]$ by the induction hypothesis.

Now note that, by the hint, we have $\sum_{i=l}^r A[i] = 0$ if and only if $\sum_{i=1}^{l-1} A[i] = \sum_{i=1}^r A[i]$ or in other words, $B[l - 1] = B[r]$. Thus, the algorithm only needs to check if there are two indices $1 \leq i < j \leq n$, where $B[i] = B[j]$. We prove the second part of the algorithm does that.

In the algorithm $C[i]$ will contain the number of elements in array B which have value $i - 5n - 1$. This follows from proof of correctness of counting sort. If the frequency of any element is greater than 1, this means that two different indices in B have the same value and thus the answer should be *Yes*; otherwise the answer is *No*: this is exactly what is done by the algorithm, proving the correctness.

Runtime Analysis: Creating the prefix sum array takes $O(n)$ time; creating C and running the search takes two for-loop each with $O(n)$ iteration, again taking $O(n)$ time. So the runtime is $O(n)$.

- (b) Now suppose that there is no promise on the range of the entries of A . Design a randomized algorithm for this problem with expected worst-case runtime of $O(n)$. **(10 points)**

Solution. A complete solution has three steps, algorithm, proof of correctness, and runtime analysis.

Algorithm: We again create the prefix sum array B as follows and again search if there are two indices i and j where $B[i] = B[j]$. The only difference is that since we do not have a bound on the range of the elements in A , we use hashing to find if there are duplicates in B without (implicitly) sorting B .

- (a) Create the prefix sum array B as before.
- (b) Pick a near-universal hash family and construct a hash table T of size $m = n$ using this hash function and the chaining method for handling collisions.
- (c) For $i = 1$ to n ,
 - i. If $T.\text{search}(B[i])$ is true, return *Yes*.
 - ii. Else, insert $B[i]$ to the hash table T .
- (d) Return *No*.

Proof of correctness: Suppose first that array B has a duplicate and k is the first index where there exist $j < k$ such that $B[j] = B[k]$. In this case, before inserting $B[k]$, the value $B[j]$ already exists in T (as we have inserted all previous entries of B to T), and thus we find the duplicate and return *Yes* correctly.

On the other hand, if array B has no duplicate, we will never find any $B[i]$ inside the table before inserting it and after the for-loop, we return *No* correctly.

Runtime Analysis: Creating the prefix sum array takes $O(n)$ time and the hash table all take deterministically $O(n)$ time. Each search also in *expectation* takes $O(1 + n/m) = O(1)$ time as we are using a randomized near-universal hash functions on a table of size $m = n$ and we insert at most n elements in the hash table. By linearity of expectation, the total expected runtime of the for-loop is also $O(n)$. Thus, the expected worst-case runtime of the algorithm is $O(n)$.

Problem 4. We want to purchase an item of price n and for that we have an unlimited (!) supply of three types of coins with values 5, 9, and 13, respectively. Our goal is to purchase this item using the *smallest* possible number of coins or outputting that this is simply not possible. Design a dynamic programming algorithm for this problem with worst-case runtime of $O(n)$. **(25 points)**

Example. A couple of examples for this problem:

- Given $n = 17$, the answer is “not possible” (try it!).
- Given $n = 18$, the answer is 2 coins: we pick 2 coins of value 9 (or 1 coin of value 5 and 1 of value 13).
- Given $n = 19$, the answer is 3 coins: we pick 1 coin of value 9 and 2 coins of value 5.
- Given $n = 20$, the answer is 4 coins: we pick 4 coins of value 5.
- Given $n = 21$, the answer is “not possible” (try it!).
- Given $n = 22$, the answer is 2 coins: we pick 1 coin of value 13 and 1 coin of value 9.
- Given $n = 23$, the answer is 3 coins: we pick 1 coin of value 13 and 2 coins of value 5.

Solution. We will apply the two steps for solving a dynamic programming problem: *Specification* and *Solution*. Only then, we turn our recursive formula into an algorithm (using memoization) and analyze the runtime.

Specification:

- For any integers $1 \leq i$, define:
 $K(i)$: the minimum number of coins required to have a total value of i ; if it is not possible to purchase the item using any combination of coins, we *define* $K(i) = +\infty$.

To return the answer, we simply need to return $K(n)$.

Solution:

$$K(i) = \begin{cases} +\infty & \text{if } i < 0 \\ 0 & \text{if } i = 0 \\ 1 + \min\{K(i-5), K(i-9), K(i-13)\} & \text{otherwise} \end{cases}$$

We now prove the correctness of this solution.

By definition, there is no way for us to purchase an item of negative value and so $K(i) = +\infty$ for $i < 0$ is correct (just by definition).

Let us consider the other base case of $i = 0$. In this case, by the specification, $K(0) = 0$. This is obviously correct, because if an item costs 0 (i.e., is free), we would need no coins to purchase it.

Now let us consider larger values of i . Because we know that we can only use the coins 5, 9, and 13, we know that every combination of coins that can purchase an item of value $i > 0$ must contain at least one of those three coins. If we pick coin 5 to begin, then we end up using one coin and have to purchase the remaining amount which is $i - 5$; thus, in this case, the number of coins will be $1 + K(i - 5)$ by the definition of $K(i - 5)$. Similarly, if we decide to use coin 9 or 13, then we have to pay $1 + K(i - 9)$ and $1 + K(i - 13)$, respectively. As our goal is to use the minimum number of coins, taking the minimum of these three possible options, gives us the correct answer.

Dynamic Programming Algorithm (Memoization):

We will store an array $D[1 : n]$ initialized with ‘undefined’ everywhere.

MemCoin(i):

1. if $i < 0$: return $+\infty$;
2. if $i = 0$: return 0
3. if $D[i] \neq \text{‘undefined’}$: return $D[i]$
4. Otherwise, let $D[i] = 1 + \min\{\text{MemCoin}(i - 5), \text{MemCoin}(i - 9), \text{MemCoin}(i - 13)\}$
5. return $D[i]$

This concludes the algorithm. The correctness follows from the correctness of recursive formula.

Runtime: Our memoization algorithm runs in $O(n)$ time, as there are n subproblems and each subproblem, ignoring the time it takes to do the inner recursions, takes $O(1)$ time.